


HW6

Anish Mohan

November 9, 2015

Q1

```
x=rnorm(100)
eps=rnorm(100)
y=5+2*x+7*x*x+3*x*x*x+5*eps
plot(x,y)
```




```
u=lm(y~-1)
u0=lm(y~1)
u1=lm(y~poly(x,1))
u2=lm(y~poly(x,2))
u3=lm(y~poly(x,3))
u4=lm(y~poly(x,4))
u5=lm(y~poly(x,5))
xlim=range(x)
xgrid=seq(from=xlim[1],to=xlim[2], length.out = 100)
```

- 1a. $\lambda = \infty, m=0$

Since $\lambda = \infty$, to minimize the equation, we need $g^{(0)}$ be zero or tend towards zero. This is possible with a function $g(x)=0$


```
plot(x,y)
tempval=rep(0,100)
lines(xgrid,tempval)
```



- 1b. $\lambda = \infty, m=1$

Since $\lambda = \infty$, to minimize the equation, we need $g^{(1)}$ to be lowest or tend towards zero. This is only possible if $g(x)$ is a constant function i.e $g(x) = \beta_0$


```
plot(x,y)
val=u0$coefficients[1]
tempval=rep(val,100)
lines(xgrid,tempval)
```



- 1c. $\lambda = \infty$, $m=2$

Since $\lambda = \infty$, to minimize the equation, we need $g^{(2)}$ to be lowest or tend towards zero. This is possible if $g^{(1)}$ is constant i.e $g(x)$ is a linear function with constant slope. Hence the function would be $g(x) = \beta_0 + \beta_1 * x$.


```
plot(x,y)
xlim=range(x)
xgrid=seq(from=xlim[1],to=xlim[2], length.out = 100)
pred=predict(u1,newdata=list(x=xgrid),se=T)
lines(xgrid, pred$fit, lwd=2)
```



- 1d. $\lambda = \infty$, $m=3$

Since $\lambda = \infty$, to minimize the equation, we need $g^{(3)}$ to be lowest or tend towards zero. This is possible if $g^{(2)}$ is constant. This implied $g^{(1)}$ is uniformly increasing or has constant slope i.e the function $g(x)$ would be increasing at a constant rate. This is possible with a function $g(x) = \beta_0 + \beta_1 * x + \beta_2 * x^2$


```
plot(x,y)
xlim=range(x)
xgrid=seq(from=xlim[1],to=xlim[2], length.out = 100)
pred=predict(u2,newdata=list(x=xgrid),se=T)
lines(xgrid, pred$fit, lwd=2)
```



- 1e. $\lambda = 0, m=3$

Since $\lambda = 0$, to minimize the equation, we need residual error to be lowest. This is possible with a function that minimizes the RSS which in this case is $g(x) = \beta_0 + \beta_1 * x + \beta_2 * x^2 + \beta_3 * x^3$ as shown by low p value from the summary results.

```
plot(x,y)
xlim=range(x)
xgrid=seq(from=xlim[1],to=xlim[2], length.out = 100)
pred=predict(u3,newdata=list(x=xgrid),se=T)
lines(xgrid, pred$fit, lwd=2)
```



```
summary(u4)
```

```
##
## Call:
## lm(formula = y ~ poly(x, 4))
##
## Residuals:
```

	Min	1Q	Median	3Q	Max
	-15.6386	-3.1319	0.6367	3.8068	9.1043

```
##
## Coefficients:
```


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	9.492	0.517	18.362	< 2e-16 ***
poly(x, 4)1	48.803	5.170	9.440	2.60e-15 ***
poly(x, 4)2	39.694	5.170	7.678	1.42e-11 ***
poly(x, 4)3	43.296	5.170	8.375	4.85e-13 ***
poly(x, 4)4	6.683	5.170	1.293	0.199

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 5.17 on 95 degrees of freedom
## Multiple R-squared:  0.6983, Adjusted R-squared:  0.6856
## F-statistic: 54.97 on 4 and 95 DF, p-value: < 2.2e-16
```

Q2

```
x=-2:2
beta0=1
beta1=1
beta2=-2
y=rep(NA,5)

for(i in 1:5 ) {
  if(x[i]<1){
    y[i]=beta0+x[i]*beta1
  } else {
    y[i]=beta0+beta1*x[i]+beta2*(x[i]-1)^2
  }
}
plot(x,y)
```



Q3

- 3a. Functions g_2^{λ} will have a smaller training error as it is higher order polynomial and has higher order penalty function. It will fit the training data well
- 3b. Functions g_1^{λ} could have a smaller test error as the higher order function g_2^{λ} will tend to overfit the noise.
- 3c. Functions g_1^{λ} and g_2^{λ} are same when $\lambda = 0$, hence they will have same training and test error.

Q4

```
library(ISLR)

## Warning: package 'ISLR' was built under R version 3.2.2

library(boot)

## Warning: package 'boot' was built under R version 3.2.2

attach(Auto)

set.seed(1)
#10-Fold validation comparing a linear and cubic model


fit.lm=glm(mpg~horsepower, data=Auto)
cv.error=cv.glm(Auto,fit.lm,k=10)
cv.error$delta

## [1] 24.10716 24.09865

hplims=range(Auto$horsepower)
hp.grid=seq(from=hplims[1],to=hplims[2])

plot(horsepower, mpg, xlim=hplims, cex=0.5, col="darkgrey")
preds=predict(fit.lm,newdata=list(horsepower=hp.grid), se=T)
lines(hp.grid, preds$fit, lwd=2, col="blue")
title("Applying linear regression")
```

Applying linear regression



```

fit.poly=glm(mpg~poly(horsepower,3),data=Auto)
cv.error=cv.glm(Auto,fit.poly,K=10)
cv.error$delta

## [1] 19.37347 19.35072

print(paste0("Error from polynomial regression is smaller than linear
regression" ))

## [1] "Error from polynomial regression is smaller than linear regression"


#Using Polynomial regression

preds=predict(fit.poly,newdata=list(horsepower=hp.grid), se=T)
se.bands=cbind(preds$fit+2*preds$se,preds$fit-2*preds$se.fit)

plot(horsepower, mpg, xlim=hplims, cex=0.5, col="darkgrey")
title("Applying polynomial regression")
lines(hp.grid, preds$fit, lwd=2, col="blue")
matlines(hp.grid, se.bands, lwd=1, col="blue", lty=3)

```

Applying polynomial regression



```

#using ANOVA to analyze variance.
fit.1=lm(mpg~horsepower, data=Auto)
fit.2=lm(mpg~poly(horsepower,2), data=Auto)
fit.3=lm(mpg~poly(horsepower,3), data=Auto)
fit.4=lm(mpg~poly(horsepower,4), data=Auto)

anova(fit.1,fit.2,fit.3,fit.4)

```


```
## Analysis of Variance Table
##
## Model 1: mpg ~ horsepower
## Model 2: mpg ~ poly(horsepower, 2)
## Model 3: mpg ~ poly(horsepower, 3)
## Model 4: mpg ~ poly(horsepower, 4)
##   Res.Df    RSS Df Sum of Sq      F Pr(>F)
## 1     390 9385.9
## 2     389 7442.0   1   1943.89 101.6666 <2e-16 ***
## 3     388 7426.4   1    15.59   0.8155 0.3670
## 4     387 7399.5   1    26.91   1.4076 0.2362
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

print(paste0("Quadratic model seems to be a good fit"))

## [1] "Quadratic model seems to be a good fit"


#Applying spline model
library(splines)
fit=lm(mpg~bs(horsepower, df=6),data=Auto)
pred=predict(fit,newdata=list(horsepower=hp.grid),se=T)
plot(horsepower,mpg, col="gray")
title('Applying splines')
lines(hp.grid, pred$fit,lwd=2)
lines(hp.grid, pred$fit+2*pred$se,lty="dashed")
lines(hp.grid, pred$fit-2*pred$se,lty="dashed")
```

Applying splines



```
#Applying Local regression
fit=loess(mpg~horsepower, span=0.2, data=Auto)
pred=predict(fit, newdata=data.frame(horsepower=hp.grid))
plot(horsepower,mpg, col="gray")
title('Applying Local regression')
lines(hp.grid, pred,col="blue", lwd=2)
```

Applying Local regression



Q5

- 5a.

```
library(MASS)
attach(Boston)

fit.poly=glm(nox~poly(dis,3))
print(paste0("Coefficient of the polynomial eqn"))

## [1] "Coefficient of the polynomial eqn"

coef(fit.poly)

## (Intercept) poly(dis, 3)1 poly(dis, 3)2 poly(dis, 3)3
## 0.5546951 -2.0030959 0.8563300 -0.3180490


summary(fit.poly)
```

```
##
## Call:
## glm(formula = nox ~ poly(dis, 3))
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -0.121130  -0.040619  -0.009738   0.023385   0.194904
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   0.554695   0.002759  201.021 < 2e-16 ***
## poly(dis, 3)1 -2.003096   0.062071  -32.271 < 2e-16 ***
## poly(dis, 3)2  0.856330   0.062071   13.796 < 2e-16 ***
## poly(dis, 3)3 -0.318049   0.062071   -5.124 4.27e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for gaussian family taken to be 0.003852802)
##
##      Null deviance: 6.7810  on 505  degrees of freedom
## Residual deviance: 1.9341  on 502  degrees of freedom
## AIC: -1370.9
##
## Number of Fisher Scoring iterations: 2

dislims=range(Boston$dis)
dis.grid=seq(from=dislims[1],to=dislims[2])
preds=predict(fit.poly,newdata=list(dis=dis.grid), se=T)
se.bands=cbind(preds$fit+2*preds$se,preds$fit-2*preds$se.fit)


plot(dis,nox, xlim=dislims,xlab="Distance from Employment Ctr",
ylab="Notrogen Oxides Conc", cex=0.5, col="darkgrey")
title("Polynomial Regression: Boston Dataset")
lines(dis.grid, preds$fit, lwd=2, col="blue")
matlines(dis.grid, se.bands, lwd=1, col="blue", lty=3)
```

Polynomial Regression: Boston Dataset



- 5b.

```
fit=matrix(0,nrow=10,ncol=1)
plot(dis,nox, xlim=dislims,xlab="Distance from Employment Ctr",
      ylab="Nitrogen Oxides Conc", cex=0.5, col="darkgrey")
for (i in 1:10){
  lm.fit=glm(nox~poly(dis,i),data=Boston)
  fit[i]=sum(lm.fit$residuals^2)
  preds=predict(lm.fit,newdata=list(dis=dis.grid), se=T)
  lines(dis.grid, preds$fit, lwd=2, col=i)
}
```




```
fit
```

```
##           [,1]
## [1,] 2.768563
## [2,] 2.035262
## [3,] 1.934107
## [4,] 1.932981
## [5,] 1.915290
## [6,] 1.878257
## [7,] 1.849484
## [8,] 1.835630
## [9,] 1.833331
## [10,] 1.832171
```

- 5c.

```
set.seed(1)
cv.error.delta=rep(NA,10)
for(i in 1:10){
  fit.poly=glm(nox~poly(dis,i),data=Boston)
  cv.error=cv.glm(Boston,fit.poly,k=10)
  cv.error.delta[i]=cv.error$delta[1]
}
plot(cv.error.delta)
```



- Degree 4 polynomial is the simplest model with lowest Cross-Validation error.
- A linear model does not fit the data, hence it has higher error. The quadratic and cubic models are comparable as the difference in error between them on a 10 fold CV set is small. Degree 4 polynomial shows the smallest error.
- For polynomials > 4, the residual errors are high implying that the model fits to the training data but gives poor results on the CV test data.
- 5d.


```
library(splines)
fit=lm(nox~bs(dis,df=4),data=Boston)
dislims=range(Boston$dis)
dis.grid=seq(from=dislims[1],to=dislims[2])
summary(fit)

##
## Call:
## lm(formula = nox ~ bs(dis, df = 4), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.124622 -0.039259 -0.008514  0.020850  0.193891
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
```



```
## (Intercept)      0.73447    0.01460  50.306 < 2e-16 ***
## bs(dis, df = 4)1 -0.05810    0.02186  -2.658  0.00812 **
## bs(dis, df = 4)2 -0.46356    0.02366 -19.596 < 2e-16 ***
## bs(dis, df = 4)3 -0.19979    0.04311  -4.634  4.58e-06 ***
## bs(dis, df = 4)4 -0.38881    0.04551  -8.544 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06195 on 501 degrees of freedom
## Multiple R-squared:  0.7164, Adjusted R-squared:  0.7142
## F-statistic: 316.5 on 4 and 501 DF,  p-value: < 2.2e-16

pred=predict(fit,newdata=list(dis=dis.grid), se=T)
plot(dis,nox, col="gray")
lines(dis.grid, pred$fit, lwd=2)
lines(dis.grid, pred$fit+2*pred$se, lty="dashed")
lines(dis.grid, pred$fit-2*pred$se, lty="dashed")
```




- Knots were chosen by specifying the df attribute that chose the knots at uniform quantile of data. In this case, the knots were at 25th, 50th and 75th quantiles
- 5e.

```
plot(dis,nox, col="gray", xlab="Degrees of Freedom")
rss=rep(NA,10)
for(i in 3:10){
  lm.fit=lm(nox~bs(dis,df=i),data=Boston)
```


```

pred=predict(lm.fit,newdata=list(dis=dis.grid), se=T)
lines(dis.grid, pred$fit, lwd=2, col=i*10)
rss[i]=sum(lm.fit$residuals^2)
}

```




```
plot(rss)
```



- RSS decreases monotonically as we increase degrees of freedom. Lowest RSS is with $df=10$
- 5f

```
set.seed(1)
cv.error.delta=rep(NA,10)
for(i in 1:10){
  fit.spline=glm(nox~bs(dis,df=i),data=Boston)
  cv.error=cv.glm(Boston,fit.spline,k=10)
  cv.error.delta[i]=cv.error$delta[1]
}
plot(cv.error.delta, xlab="Degrees of Freedom")
```



- Lowest error obtained through cross validation is for $df=10$.
- As we increase the degrees of freedom, initially, till we reach $df=3$, there is increase in the RSS. However, error starts decreasing after that till $df=5$. Errors fluctuate a bit, before getting the lowest error at $df=10$.

Q6

```
x1=rnorm(100)
x2=rnorm(100)
eps=rnorm(100)
y=5+2*x1+7*x2+eps

beta0=rep(NA,1000)
beta1=rep(NA,1000)
beta2=rep(NA,1000)

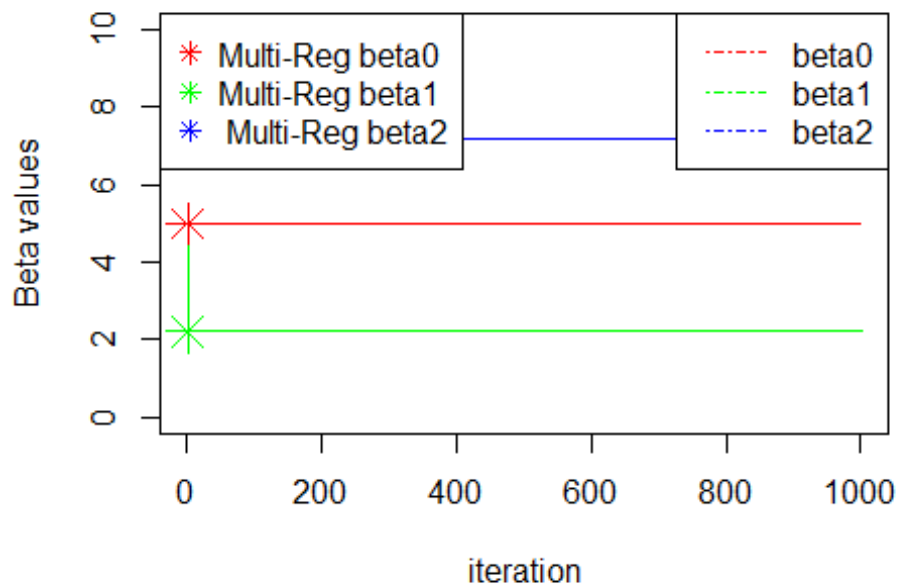
beta1[1]=5

for(i in 1:1000){
  a=y-beta1[i]*x1
  beta2[i]=lm(a~x2)$coef[2]
  beta0[i]=lm(a~x2)$coef[1]

  a=y-beta2[i]*x2
  beta1[i+1]=lm(a~x1)$coef[2]
}


plot(1:1001,beta1, col="green", type='l', ylim=c(0,10),
xlab="iteration", ylab="Beta values")
lines(1:1000,beta0, col="red")
lines(1:1000,beta2, col="blue")

u=lm(y~x1+x2)
points(coef(u)[1],col="red", pch=8, cex=2)
points(coef(u)[2],col="green", pch=8, cex=2)
points(coef(u)[3],col="blue", pch=8, cex=2)
legend("topright",c("beta0", "beta1", "beta2"), lty=4,col =
c("red", "green", "blue"))
legend("topleft",c("Multi-Reg beta0", "Multi-Reg beta1", " Multi-Reg
beta2"), pch=8,col = c("red", "green", "blue"))
```



```
plot(1:1001,beta1, col="green", type='l', ylim=c(0,10),
xlab="iteration", ylab="Beta values")
title("Showing abline from multiple regression coefficients.")
lines(1:1000,beta0, col="red")
lines(1:1000,beta2, col="blue")
abline(h=beta0, lty="dashed")
abline(h=beta1, lty="dashed")
abline(h=beta2, lty="dashed")
```

Showing abline from multiple regression coefficient



```
print(paste0("Within 2nd approximations the backfitting answers  
were good estimates"))
```

```
## [1] "Within 2nd approximations the backfitting answers were good  
estimates"
```

```
beta0[1:5]
```

```
## [1] 5.149512 5.014518 5.009577 5.009396 5.009389
```

```
beta1[1:5]
```

```
## [1] 5.000000 2.319657 2.221546 2.217955 2.217824
```

```
beta2[1:5]
```

```
## [1] 6.693492 7.145315 7.161853 7.162459 7.162481
```