

HW7

Anish Mohan

November 16, 2015

Q1

Q2

- Bootstrap estimates for $P(\text{Class is Red} \mid X) = \{0.1, 0.15, 0.2, 0.2, 0.55, 0.6, 0.6, 0.65, 0.7, 0.75\}$
- Majority Approach: $X = \text{Red}$
- Number of times $P(\text{Class is Red} \mid X) > 0.5$. In this case 6 times $P(\text{Class is Red} \mid X) > 0.5$ hence $X = \text{Red}$
- Average Approach: $X = \text{Green}$
- Take average of the probability values. i.e $P(\text{Class is Red} \mid X) = 0.45$. Hence $P(\text{Class is Green} \mid X) = 0.55$.
 $X = \text{Green}$

Q3

- 3a.

```
library(tree)
library(ISLR)
attach(Carseats)
set.seed(1)
train=sample(1:nrow(Carseats),200)
Carseats.train=Carseats[train,]
Carseats.test=Carseats[-train,]
```

- 3b.

```
tree.carseats=tree(Sales~.,Carseats,subset=train)
summary(tree.carseats)
```


```
##
## Regression tree:
## tree(formula = Sales ~ ., data = Carseats, subset = train)
## Variables actually used in tree construction:
## [1] "ShelveLoc" "Price" "Age" "Advertising" "Income"
## [6] "CompPrice"
## Number of terminal nodes: 18
## Residual mean deviance: 2.36 = 429.5 / 182
## Distribution of residuals:
## Min. 1st Qu. Median Mean 3rd Qu. Max.
## -4.2570 -1.0360 0.1024 0.0000 0.9301 3.9130
```

```
par(mfrow=c(1,1))

tree.pred=predict(tree.carseats,Carseats.test)
mean((tree.pred-Carseats.test$Sales)^2)
```


```
## [1] 4.148897
```

```
plot(tree.carseats)
text(tree.carseats)
```




- Variables that were included in the construction of the tree were: ShelveLoc, Price, Age, Advertising, Income and Comp Price. There are 18 terminal nodes or leaves. The RSS for the training data is 2.36
 - The key variable is ShelveLoc which is split into 2 parts Bad and Medium. ShelveLoc = Good is not a critical indicator of sale price. Price is another key variable which is combined when with ShelveLoc=Bad to give to give split at 120.5 and combined with ShelveLoc=Medium to give a split at price=113. Best Sales number seems to be when the ShelveLoc is a Medium and the sale price is 113.
- the test MSE is 4.15
- 3c.

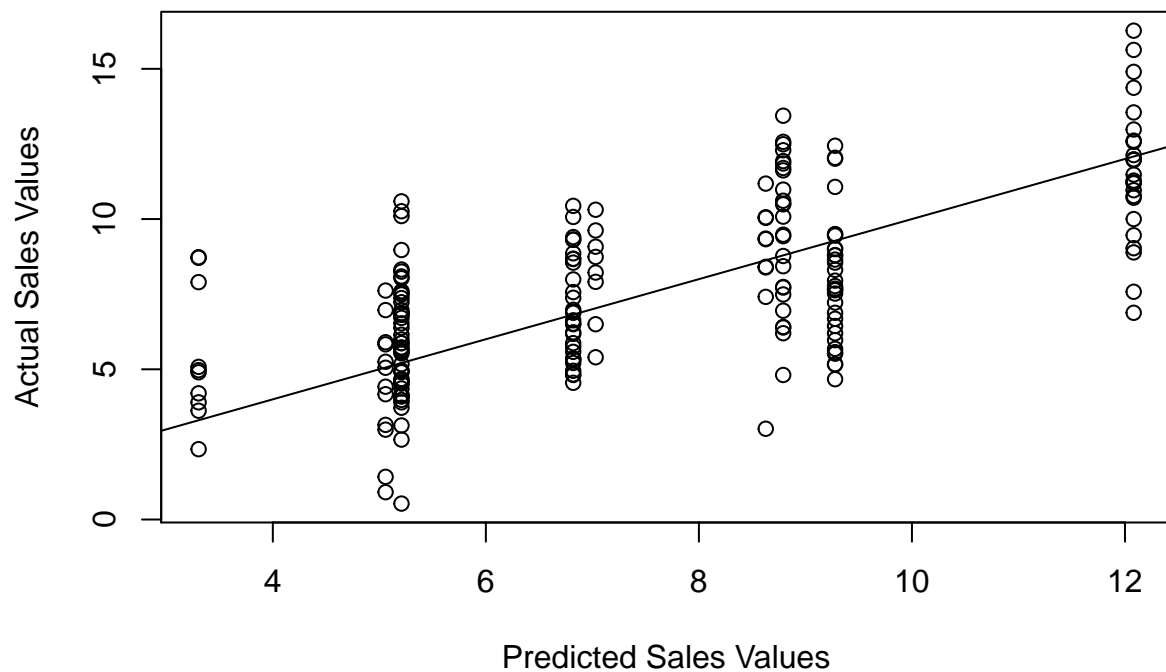
```
cv.carseats=cv.tree(tree.carseats)
plot(cv.carseats$size,cv.carseats$dev, type='b')
```



```
prune.carseats=prune.tree(tree.carseats,best=9)
plot(prune.carseats)
text(prune.carseats, pretty=0)
```



```
yhat=predict(prune.carseats, newdata=Carseats.test)
plot(yhat, Carseats.test$Sales, xlab="Predicted Sales Values", ylab="Actual Sales Values")
abline(0,1)
```



```
mean((yhat-Carseats.test$Sales)^2)
```

```
## [1] 4.993124
```

- Pruning the tree does not reduce the test MSE. Infact the test MSE increases to 4.99
- 3d.

```
library(randomForest)
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```


```
set.seed(1)
```

```
bag.carseats=randomForest(Sales~.,data=Carseats,subset=train,mtry=10,importance=T)
```

```
yhat.bag=predict(bag.carseats,newdata=Carseats.test)
```

```
plot(yhat.bag, Carseats.test$Sales)
```

```
abline(0,1)
```




```
mean((yhat.bag-Carseats.test$Sales)^2)
```

```
## [1] 2.554292
```


```
varImpPlot(bag.carseats)
```

bag.carseats



- Test MSE after doing bagging is 2.55
- Price and ShelfLoc are the two most important variables for determining amount of Sales
- 3e.

```
set.seed(1)
rf.carseats=randomForest(Sales~.,data=Carseats,subset=train,importance=T)
yhat.rf=predict(rf.carseats,newdata=Carseats.test)
plot(yhat.rf, Carseats.test$Sales)
abline(0,1)
```




```
mean((yhat.rf-Carseats.test$Sales)^2)
```

```
## [1] 3.30763
```

```
varImpPlot(rf.carseats)
```

rf.carseats



- Test MSE after doing bagging is 3.3 which is higher than bagging but lower than regression tree
- Price and ShelfLoc are the two most important variables for determining amount of Sales
- We used the default 'm' for random forest that is \sqrt{p} i.e = 3 variables. Just using 3 variables increased the Test MSE compared to Bagging which used 10 variables.

Q4

- 4a.

```
library(ISLR)
fix(Hitters)
Hitters=na.omit(Hitters)
Hitter=na.omit(Hitters)
Hitter$Salary=log(Hitters$Salary)
```

- 4b.

```
train=1:200
test=-train
Hitter.train=Hitter[train,]
Hitter.test=Hitter[-train,]
```

- 4c.

```
library(gbm)
```

```
## Loading required package: survival
## Loading required package: splines
## Loading required package: lattice
## Loading required package: parallel
## Loaded gbm 2.1.1
```


```
set.seed(1)

train.mse=matrix(data=NA, nrow=100, ncol=2)
test.mse=matrix(data=NA, nrow=100, ncol=2)
for(i in 1:100){
  lambda=i*0.01
  train.mse[i,1]=lambda
  test.mse[i,1]=lambda

  boost.hitter=gbm(Salary~.,data=Hitter.train, distribution="gaussian", n.trees=1000,interaction.d
  train.mse[i,2]=mean(boost.hitter$train.error)


  yhat.boost=predict(boost.hitter,newdata=Hitter.test,n.trees=1000)
  test.mse[i,2]=mean((yhat.boost-Hitter.test$Salary)^2)
}

plot(train.mse[,1],train.mse[,2], xlab="Shrinkage factor", ylab="Training MSE")
```



- 4d.

```
plot(test.mse[,1],test.mse[,2], xlab="Shrinkage factor", ylab="Test MSE")
```

- 4e.

```
lm.fit=lm(Salary~.,data=Hitter.train)
lm.pred=predict(lm.fit,Hitter.test)
mean((lm.pred-Hitter.test$Salary)^2)
```

```
## [1] 0.4917959
```

```
library(glmnet)
```

```
## Loading required package: Matrix
## Loading required package: foreach
## Loaded glmnet 2.0-2
```


```
set.seed(1)
x=model.matrix(Salary~.,Hitter)
y=Hitter$Salary
y.test=y[test]
grid=10^seq(10,-2,length=100)
ridge.mod=glmnet(x[train,],y[train],alpha=0, lambda=grid, thresh=1e-12)
ridge.pred=predict(ridge.mod,s=212,newx=x[test,])
mean((ridge.pred-y.test)^2)
```

```
## [1] 0.6313728
```

- Minimum Test MSE: Boosting: 0.249
- Minimum Test MSE: Linear Model fit: 0.49
- Minimum Test MSE: Ridge Regression: 0.45

- 4f.

```
boost.hitter=gbm(Salary~.,data=Hitter.train, distribution="gaussian", n.trees=1000,interaction.depth=4)
summary(boost.hitter)
```




```
##          var      rel.inf
## CatBat    CatBat 21.2500123
## CRuns     CRuns 14.1765669
## CWalks    CWalks 8.7472089
## CRBI      CRBI  7.1509000
## PutOuts   PutOuts 6.7109573
## Walks     Walks  6.3698013
## AtBat     AtBat  5.2080574
## CHmRun    CHmRun 4.6710617
## Years     Years  4.5471386
## RBI       RBI   4.2839575
## Assists   Assists 3.9848415
## Hits      Hits  2.9274869
## Runs      Runs  2.6974009
## HmRun     HmRun  2.4462063
## Errors    Errors 1.9190240
## CHits     CHits  1.0183290
## NewLeague NewLeague 0.8562599
## Division  Division 0.7858884
## League    League  0.2489012
```

– CatBat and CRuns are the key variables.

- 4g.

```
library(randomForest)
set.seed(1)
bag.Hitter=randomForest(Salary~.,data=Hitter.train,mtry=19,importance=T)
yhat.bag=predict(bag.Hitter,newdata=Hitter.test)
```

```
plot(yhat.bag, Hitter.test$Salary)
abline(0,1)
```



```
mean((yhat.bag-Hitter.test$Salary)^2)
```

```
## [1] 0.228722
```

– Test MSE using Bagging: 0.23

Q5

```
```\r
library(kernlab)
library(e1071)
set.seed(1)
x=matrix(rnorm(100*2),ncol=2)
x[1:25,]=x[1:25,]+2
x[26:50,]=x[26:50,]-2

y=c(rep(1,60),rep(2,40))
df=data.frame(x=x, y=as.factor(y))

plot(x,col=y)
```\r

![] (HW7_files/figure-latex/unnamed-chunk-14-1.pdf)

```\r
train=sample(100,75)
```

```

test=-train

df.train=df[train,]
df.test=df[test,]

svmfit=svm(y~.,data=df[train,],kernel="radial", gamma=1,cost=1,scale=F)
plot(svmfit,df[train,])
```



```r
summary(svmfit)
```



```

##
Call:
svm(formula = y ~ ., data = df[train,], kernel = "radial", gamma = 1,
cost = 1, scale = F)
##
Parameters:
SVM-Type: C-classification
SVM-Kernel: radial
cost: 1
gamma: 1
##
Number of Support Vectors: 52
##
(25 27)
##
##
Number of Classes: 2
##
Levels:
1 2
```

```r
ytrain=predict(svmfit, df[train,])
table(predict=ytrain, truth=df[train,"y"])
```



```

##
truth
predict 1 2
1 36 0
2 10 29
```

```r
ytest=predict(svmfit, df[test,])
table(predict=ytest, truth=df[test,"y"])

```


```


```

```

...

...

truth
predict 1 2
1 11 1
2 3 10
...

```r
#Linear filter
svmfit=svm(y~.,data=df[train,],kernel="linear", cost=0.1,scale=F)
plot(svmfit,df[train,])
...

![] (HW7_files/figure-latex/unnamed-chunk-14-3.pdf)

```r
summary(svmfit)
...

...

##
Call:
svm(formula = y ~ ., data = df[train,], kernel = "linear", cost = 0.1,
scale = F)
##
##
Parameters:
SVM-Type: C-classification
SVM-Kernel: linear
cost: 0.1
gamma: 0.5
##
Number of Support Vectors: 60
##
(29 31)
##
##
Number of Classes: 2
##
Levels:
1 2
...

```r
ytrain=predict(svmfit, df[train,])
table(predict=ytrain, truth=df[train,"y"])
...

...

##          truth
## predict  1  2
##          1 46 29

```

```
##      2  0  0
```

```r
ytest=predict(svmfit, df[test,])
table(predict=ytest, truth=df[test,"y"])
```



```
##      truth
## predict  1  2
##      1 14 11
##      2  0  0
```
```


```

- With the Radial Kernel, 65 out of 75 points have been assigned the correct class in the training data and 21 out of 25 points have been assigned the correct class in the test data.
- With the Linear kernel, only 46 out of 75 points are correctly assigned the correct class in training data and 14 out of 25 points have been assigned the correct class in the test data.

Q6

```
library(ISLR)
median(Auto$mpg)
```

```
## [1] 22.75
```

```
mileage=ifelse(Auto$mpg>=median(Auto$mpg),1,0)

df=data.frame(cylinder=Auto$cylinders,displacement= Auto$displacement, horsepower= Auto$horsepower,

set.seed(1)
tune.out=tune(svm,mileage~.,data=df, kernel="linear",ranges=list(cost=c(0.001,0.01,0.1,1,5,10,100)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.09179487
##
## - Detailed performance results:
##   cost      error dispersion
## 1 1e-03 0.13288462 0.05551120
## 2 1e-02 0.09179487 0.05812971
```

```
## 3 1e-01 0.09692308 0.06369443
## 4 1e+00 0.09179487 0.04543280
## 5 5e+00 0.10198718 0.04338864
## 6 1e+01 0.11480769 0.05828005
## 7 1e+02 0.11730769 0.06521821
```

```
bestmod=tune.out$best.model
summary(bestmod)
```

```
##
## Call:
## best.tune(method = svm, train.x = mileage ~ ., data = df, ranges = list(cost = c(0.001,
## 0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##       cost:  1
##       gamma: 0.003215434
##
## Number of Support Vectors: 118
##
## ( 52 66 )
##
##
## Number of Classes: 2
##
## Levels:
## 0 1
```

```
#radial
set.seed(1)
tune.out=tune(svm,mileage~.,data=df, kernel="radial",ranges=list(cost=c(0.001,0.01,0.1,1,5,10,100)),
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##   1      1
##
## - best performance: 0.07647436
##
## - Detailed performance results:
##   cost gamma    error dispersion
## 1 1e-03    0.5 0.56115385 0.04344202
## 2 1e-02    0.5 0.56115385 0.04344202
## 3 1e-01    0.5 0.08923077 0.05559893
## 4 1e+00    0.5 0.07897436 0.05443042
```

```
## 5 5e+00 0.5 0.08147436 0.06565669
## 6 1e+01 0.5 0.08910256 0.06377480
## 7 1e+02 0.5 0.08910256 0.06377480
## 8 1e-03 1.0 0.56115385 0.04344202
## 9 1e-02 1.0 0.56115385 0.04344202
## 10 1e-01 1.0 0.56115385 0.04344202
## 11 1e+00 1.0 0.07647436 0.05657355
## 12 5e+00 1.0 0.08160256 0.06250579
## 13 1e+01 1.0 0.08416667 0.06164376
## 14 1e+02 1.0 0.08416667 0.06164376
## 15 1e-03 2.0 0.56115385 0.04344202
## 16 1e-02 2.0 0.56115385 0.04344202
## 17 1e-01 2.0 0.56115385 0.04344202
## 18 1e+00 2.0 0.11724359 0.04962645
## 19 5e+00 2.0 0.11467949 0.05251931
## 20 1e+01 2.0 0.11467949 0.05251931
## 21 1e+02 2.0 0.11467949 0.05251931
## 22 1e-03 3.0 0.56115385 0.04344202
## 23 1e-02 3.0 0.56115385 0.04344202
## 24 1e-01 3.0 0.56115385 0.04344202
## 25 1e+00 3.0 0.33410256 0.15766416
## 26 5e+00 3.0 0.29570513 0.13702698
## 27 1e+01 3.0 0.29570513 0.13702698
## 28 1e+02 3.0 0.29570513 0.13702698
## 29 1e-03 4.0 0.56115385 0.04344202
## 30 1e-02 4.0 0.56115385 0.04344202
## 31 1e-01 4.0 0.56115385 0.04344202
## 32 1e+00 4.0 0.45391026 0.09910221
## 33 5e+00 4.0 0.46160256 0.07848745
## 34 1e+01 4.0 0.46160256 0.07848745
## 35 1e+02 4.0 0.46160256 0.07848745
```

```
bestmod=tune.out$best.model
summary(bestmod)
```

```
##
## Call:
## best.tune(method = svm, train.x = mileage ~ ., data = df, ranges = list(cost = c(0.001,
## 0.01, 0.1, 1, 5, 10, 100), gamma = c(0.5, 1, 2, 3, 4)), kernel = "radial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##         cost: 1
##        gamma: 1
##
## Number of Support Vectors: 376
##
## ( 187 189 )
##
##
## Number of Classes: 2
##
```



```
## Levels:
## 0 1
```

```
#polynomial
set.seed(1)
tune.out=tune(svm,mileage~.,data=df, kernel="polynomial",ranges=list(cost=c(0.001,0.01,0.1,1,5,10,100),gamma=c(0.001,0.01,0.1,1,5,10,100)))
summary(tune.out)
```

```
##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost gamma
##     1    0.5
##
## - best performance: 0.07634615
##
## - Detailed performance results:
##   cost gamma      error dispersion
## 1  1e-03    0.5 0.25544872 0.09514919
## 2  1e-02    0.5 0.09185897 0.06420082
## 3  1e-01    0.5 0.08160256 0.06250579
## 4  1e+00    0.5 0.07634615 0.04933672
## 5  5e+00    0.5 0.09166667 0.05108581
## 6  1e+01    0.5 0.09166667 0.04500388
## 7  1e+02    0.5 0.09179487 0.04200857
## 8  1e-03    1.0 0.09442308 0.06846018
## 9  1e-02    1.0 0.08666667 0.06734113
## 10 1e-01    1.0 0.07634615 0.04933672
## 11 1e+00    1.0 0.09166667 0.05108581
## 12 5e+00    1.0 0.09179487 0.04200857
## 13 1e+01    1.0 0.09179487 0.04200857
## 14 1e+02    1.0 0.09179487 0.04200857
## 15 1e-03    2.0 0.08923077 0.06521387
## 16 1e-02    2.0 0.07634615 0.04933672
## 17 1e-01    2.0 0.09166667 0.05108581
## 18 1e+00    2.0 0.09179487 0.04200857
## 19 5e+00    2.0 0.09179487 0.04200857
## 20 1e+01    2.0 0.09179487 0.04200857
## 21 1e+02    2.0 0.09179487 0.04200857
## 22 1e-03    3.0 0.08410256 0.05257348
## 23 1e-02    3.0 0.08660256 0.05261832
## 24 1e-01    3.0 0.09179487 0.04200857
## 25 1e+00    3.0 0.09179487 0.04200857
## 26 5e+00    3.0 0.09179487 0.04200857
## 27 1e+01    3.0 0.09179487 0.04200857
## 28 1e+02    3.0 0.09179487 0.04200857
## 29 1e-03    4.0 0.07891026 0.05003145
## 30 1e-02    4.0 0.09166667 0.05108581
## 31 1e-01    4.0 0.09179487 0.04200857
## 32 1e+00    4.0 0.09179487 0.04200857
## 33 5e+00    4.0 0.09179487 0.04200857
```

```
## 34 1e+01    4.0 0.09179487 0.04200857
## 35 1e+02    4.0 0.09179487 0.04200857
```

```
bestmod=tune.out$best.model
summary(bestmod)
```

```
##
## Call:
## best.tune(method = svm, train.x = mileage ~ ., data = df, ranges = list(cost = c(0.001,
##    0.01, 0.1, 1, 5, 10, 100), gamma = c(0.5, 1, 2, 3, 4)), kernel = "polynomial")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##     cost:    1
##   degree:    3
##   gamma:     0.5
##   coef.0:    0
##
## Number of Support Vectors:  119
##
##  ( 58 61 )
##
##
## Number of Classes:  2
##
## Levels:
##  0 1
```

Q7